

Corporate Policy Assistant

Corrected Azure AI Search Data Ingestion Process

The proposed ingestion flow is very close, but Step 4 needs one important correction: the Microsoft Foundry embedding model creates the vector embeddings. Azure AI Search does not convert those vectors a second time; it stores and indexes them for later retrieval.

Recommended Four-Step Ingestion Process

1. Azure AI Search Indexer

Reads the corporate policy documents from the Azure Blob Storage container and extracts their content. Microsoft typically refers to this component as an **indexer**, rather than an Azure AI Search batch job.

2. Azure AI Search Skillset - Text Split / Chunking

Breaks the extracted policy text into smaller chunks that can be embedded and searched efficiently. For example: Policy PDF -> extracted text -> Chunk 1, Chunk 2, Chunk 3, and so on.

3. Microsoft Foundry Embedding Model

Azure AI Search sends each policy text chunk through an embedding skill to the deployed embedding model, such as **text-embedding-3-small**. The model converts each text chunk into a numerical vector embedding.

4. Azure AI Search Index

Receives the text chunks, vector embeddings, and associated metadata and stores them in searchable index fields. Azure AI Search then builds the internal vector-search structures needed for vector, keyword, or hybrid retrieval.

Conceptual Architecture Flow

```
Azure Blob Storage
Corporate Policy Documents
|
v
Azure AI Search Indexer
Extract document content
|
v
Azure AI Search Skillset
Text Split / Chunking
|
v
Policy Text Chunks
|
v
Azure AI Search Embedding Skill
|
v
Microsoft Foundry
text-embedding-3-small
|
v
Vector Embeddings
|
v
Azure AI Search Index
- Policy text chunk
- Vector embedding
- Document ID / metadata
- Page / section information
|
v
Vector-Searchable Data
```

Important Correction to Step 4

The original Step 4 described Azure AI Search as converting vector embeddings into vector embeddings and then sending them to a separate vector database. That is not the correct model. The Foundry embedding model has already produced the vectors. Azure AI Search stores those vectors in vector fields, indexes them, and makes them searchable.

It is also clearer to think of the **Azure AI Search index** as the searchable store containing both normal text fields and vector fields, rather than treating an 'Azure AI Search Vector Index' and an 'Azure AI Search Vector Database' as two separate Azure components.

Clean Version for Architecture Documentation

Step	Component	Responsibility
1	Azure AI Search Indexer	Retrieve and extract Corporate Policy documents from Azure Blob Storage.
2	Azure AI Search Skillset	Split extracted policy text into appropriately sized text chunks.

Step	Component	Responsibility
3	Microsoft Foundry Embedding Model	Convert each policy text chunk into a numerical vector embedding.
4	Azure AI Search Index	Store and index text chunks, vector embeddings, and metadata for vector, keyword, or hybrid retrieval.