

Memory updated

Yes. When you finish the implementation, I'll help you build a consolidated **identity and RBAC matrix** for the entire Policy Assistant architecture.

It should include each identity, what component uses it, what Azure resource it needs access to, the exact role, and why that permission is required. For example:

Identity	Resource	Role	Purpose
Developer/User account	Azure AI Search	Search Index Data Reader	Test queries locally
Developer/User account	Azure AI Search	Search Service Contributor	Modify Search configuration
Function App managed identity	Azure AI Search	Search Index Data Reader	Query the policy index
Azure AI Search managed identity	Foundry/Azure OpenAI	Cognitive Services OpenAI User	Generate embeddings
Azure AI Search managed identity	Blob Storage	Storage Blob Data Reader	Read policy documents
Function App managed identity	Foundry model	Appropriate inference role	Send RAG context/query to the LLM

We can also separate the final matrix into **development permissions** versus **production/runtime permissions**, which will make it much easier to identify permissions that can be removed after implementation.

The goal should be to end with a **least-privilege security design**, not simply preserve every role that was temporarily needed while building the solution.